

Assignment 2

Logistics

You must create an IPython notebook for the problem. The notebook should clearly state the names of the group members. Please name the notebook “*HW2_Lastnames_of_group_members*”. Upload the notebook on Canvas.

The assignments are due on **March 2nd, before the beginning of class**.

Deliverables. The notebook should be split by sections, one for each question. Each section should contain the following parts:

1. A heading for each question (use a Markdown cell).
2. For each question, write a short description of your approach (another Markdown cell).
3. Write the code, and show the results or plots.

1 NYC Restaurants

We will analyze restaurant inspections in New York City, starting November 1, 2014 and ending January 31, 2015. The data is in the file `NYC_Restaurants.csv`. If you want to play with the full dataset, look at <https://data.cityofnewyork.us/Health/DOHMH-New-York-City-Restaurant-Inspection-Results/xx67-kt59>.

Setup First, read in this data using:

- `df = pd.read_csv('NYC_Restaurants.csv', dtype=unicode).`

[Q1, 7 points] Create a unique name for each restaurant. On the DataFrame created above, add a new column to your DataFrame, called 'RESTAURANT', that combines DBA, BUILDING, STREET, and ZIP-CODE fields. For example, "WENDY'S 469 FLATBUSH AVENUE 11225". Print the first 10 values of the RESTAURANT column of your DataFrame.

[Q2, 7 points] How many restaurants are included in the data? Careful now:

- A "Subway" in one particular address (i.e., building, street, zipcode, and boro) counts as one restaurant; don't combine all Subways into one restaurant!
- The data can have multiple violations for the same restaurant!

[Q3, 7 points] How many chains are there? Let us define a chain to be the same restaurant name occurring in at least two different (building, street, zipcode) addresses.

[Q4, 7 points] Plot a bar graph of the top 20 most popular chains. We already have the chains from the previous problem.

[Q5, 7 points] What fraction of all restaurants are chains?

[Q6, 7 points] Plot the number of non-chain restaurants in each boro. First, we need to figure out all the non-chain restaurants, then select out only those restaurants, and finally plot the number of such restaurants by boro. Make sure to look at the plot; we don't want to see... oh... the "missing" boro.

[Q7, 10 points] Plot the *fraction* of non-chain restaurants in each boro. The boro with the most non-chain restaurants might just be the boro with the most restaurants in general. If we want to find the boro that attracts the most "independent" restauranteurs, we must divide the number of non-chain restaurants by the total number of restaurants in the boro. Plot this.

Is the boro with the most independent restaurants also the one with the highest ratio of independent restaurants?

[Q8, 7 points] Plot the popularity of cuisines. Which cuisines are the most well-represented among all restaurants? Plot the popularity of the top 20 cuisines. Again, be careful in your counting: remember that the same restaurant can have multiple violations!

[Q9, 12 points] Plot the cuisines among restaurants which do not get cited for violations. Ideally, you should explore and see what happens when there is no violation, but here I will just tell you: the 'VIOLATION CODE' field is missing.

- Select only the 'RESTAURANT' and 'CUISINE DESCRIPTION' columns of the DataFrame (we created the 'RESTAURANT' field in an earlier question).
- Select the instances where a restaurant did *not* get cited.
- Among these instances, how frequent is each of the cuisines?

Plot the frequencies of the top-20 among these cuisines.

[Q10, 7 points] What cuisines tend to be the “cleanest”?

- Select all cuisines for which there were at least 20 inspections of restaurants representing that cuisine.
- For each such cuisine, compute the ratio of inspections that found no violations to the total number of inspections. (This is just the ratio of counts in the previous two questions.)
- Find the top-10 cuisines with the highest ratios; these are that cuisines whose restaurants are “most likely to be clean.” (*Hint: Check out the “sort_values” method for Series.*)

[Q11, 10 points] What are the most common violations in each borough? Create a table of the number of times each violation description was observed in each borough, and figure out the most common violation description for each borough.

To create the table, there are two possibilities:

- You can create a dummy column, fill it with just the number 1, and create the table using a command we have seen in class (*Hint: we will want aggfunc=sum*).
- Alternatively, you can look up the function `crosstab` which does this counting by default.

Once you do have the table, you will still need to find the most common violation description for each borough.

[Q12, 12 points] What are the most common violations per borough, *after normalizing for the relative abundance of each violation*? Hopefully, the answer to the previous question left you unsatisfied, because some violations are just very common, irrespective of borough. A better approach would be to *normalize* the violation counts, as follows.

- *Get overall frequencies:* Figure out how common each violation is, over the entire dataset; let's call this `violationFrequency`.
- *Normalize:* Consider the table of number of violations by boro that you created for the previous question. For each borough, divide the number of violations of each type by the total number of violations for that type; i.e., divide the series of violations by `violationFrequency`. We want to do this *for each borough*.
- *Find the biggest violations:* Now, after this normalization, for each borough, figure out the most common violation description.